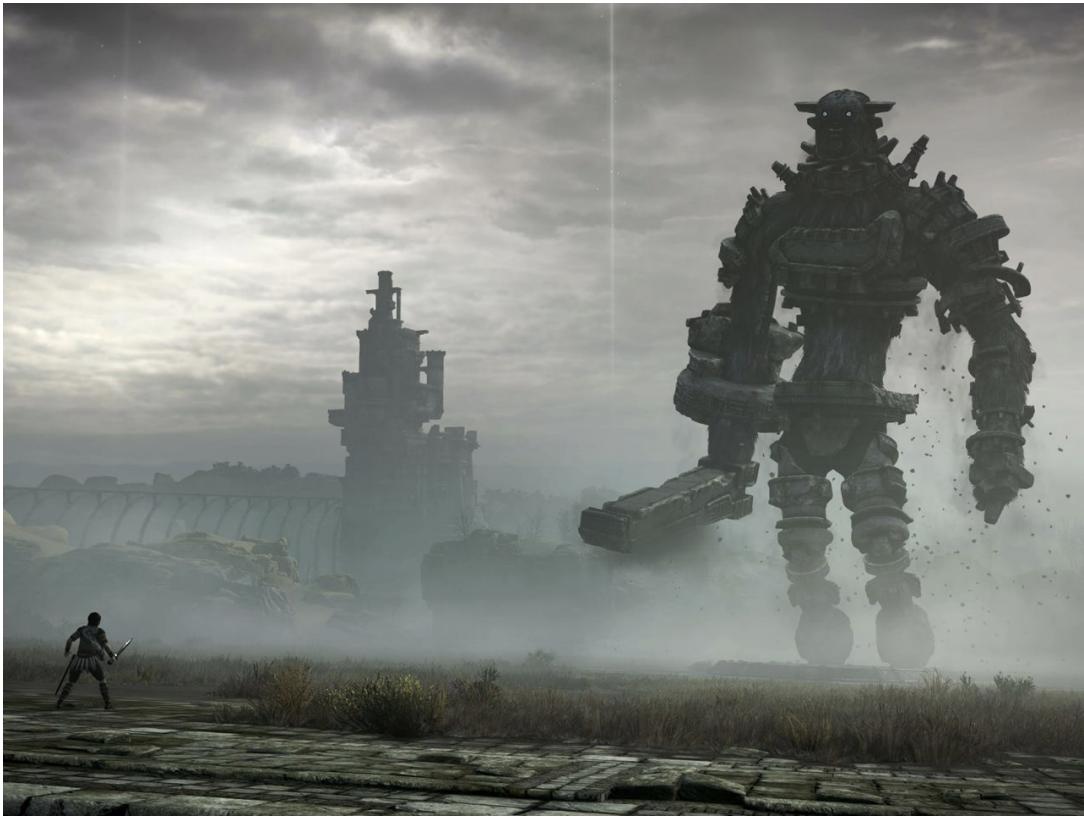


# **eth2 Red Team - Report o**

## methodology and scope



*eth2 red teaming in an image*

### **breadth-first a.k.a. trolling all over the place**

Eth2 has been under active research and development for several years before being successfully launched this year. It is likely one of the largest open-source projects of all time, driven and developed by a diverse, decentralized community. As a result, standards and best practices for red teaming and pen testing efforts on behalf of centralized organizations have only limited validity in the context of eth2.

As both relevant information and the implementation of the eth2 system is distributed across a variety of resources and as the bootstrapping of the red team research overlapped with the launch of eth2, we chose a breadth-first methodology to develop cross-domain and knowledge and identify potential weakest links in eth2's security supply chain to inform future research.

The chosen methodology also generates "real world" data on how resource and time-constrained attackers perceive and interact maliciously with the evolving eth2 ecosystem.

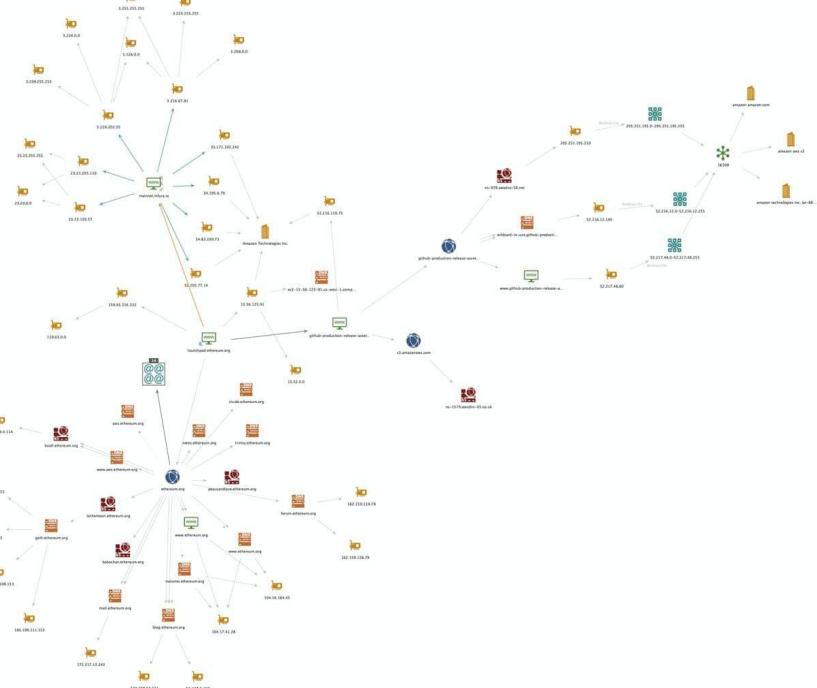
## **findings and non-findings**

[launchpad.ethereum.org](https://launchpad.ethereum.org)

## Scope

The eth2 launchpad was briefly researched and tested during the launch, as it is a high-value target for various attacks. Due to its reliance on third-party services (particularly services offered by Infura, Github, and AWS), the auxiliary services were included in the investigation, but had to be outscoped for legal reasons in the course of the analysis.

As a first step, we performed an OSINT analysis of the launchpad to understand and prepare target selection and future investigations.



From this set, we selected essentially 3 targets for a brief analysis

1. The AWS s3 bucket used by Github to host production releases of the eth2.0-deposit-CLI tool at [github-production-release-asset-2e65be.s3.amazonaws.com](https://github-production-release-asset-2e65be.s3.amazonaws.com)
  2. The Infura endpoint used to fetch deposits for the launchpad at [mainnet.infura.io/v3/9a011a6cf597453481cc94c05ba7dc54](https://mainnet.infura.io/v3/9a011a6cf597453481cc94c05ba7dc54)
  3. The Netlify instance hosting the launchpad itself at [silly-engelbart-23669a.netlify.app](https://silly-engelbart-23669a.netlify.app)

# Results

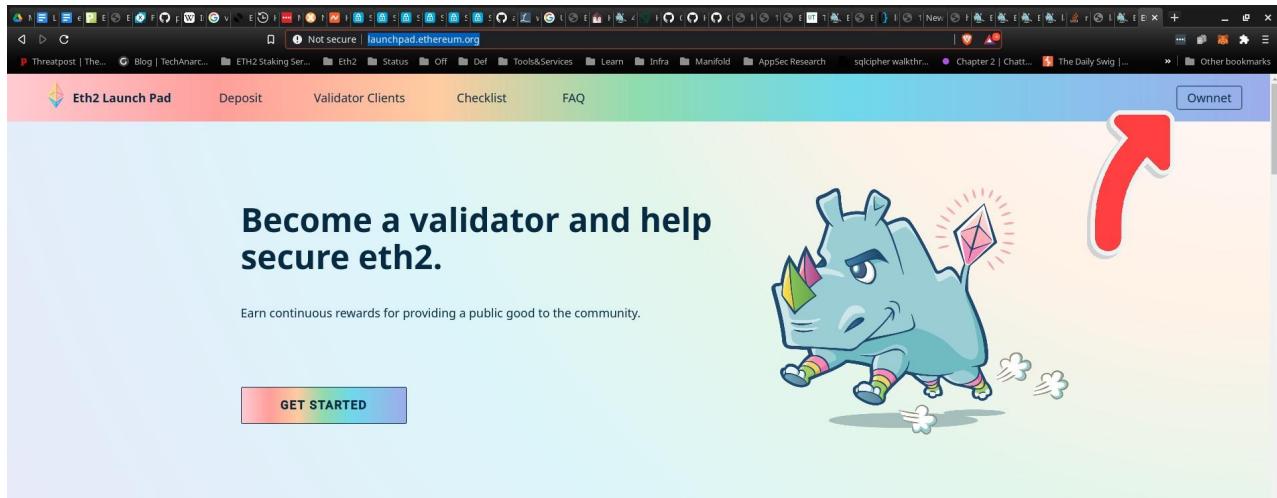
## Local Testing

Initially, we tested the launchpad and its auxiliary services in a setting where an attacker had access to the WLAN of the victim machine. In this setup, the launchpad is not safe to use and a compromise to a large extent was possible (delivering an arbitrary website with an altered deposit address or a malicious eth2-deposit-cli). Exploitable vulnerabilities are mainly but not exclusively caused by the auxiliary endpoints being mis/under configured.

|           | ssl strip            | dns spoofing | SSL security summary  |
|-----------|----------------------|--------------|---|
| AWS S3    | partially vulnerable | vulnerable   | <a href="https://www.ssllabs.com/ssltest/analyze.html?d=github-production-release-asset-2e65be.s3.amazonaws.com">https://www.ssllabs.com/ssltest/analyze.html?d=github-production-release-asset-2e65be.s3.amazonaws.com</a> |
| Infura    | partially vulnerable | vulnerable   | <a href="https://www.ssllabs.com/ssltest/analyze.html?d=mainnet.infura.io">https://www.ssllabs.com/ssltest/analyze.html?d=mainnet.infura.io</a>   |
| Launchpad | not vulnerable       | vulnerable   | <a href="https://www.ssllabs.com/ssltest/analyze.html?d=silly-engelbart-23669a.netlify.app">https://www.ssllabs.com/ssltest/analyze.html?d=silly-engelbart-23669a.netlify.app</a>   |

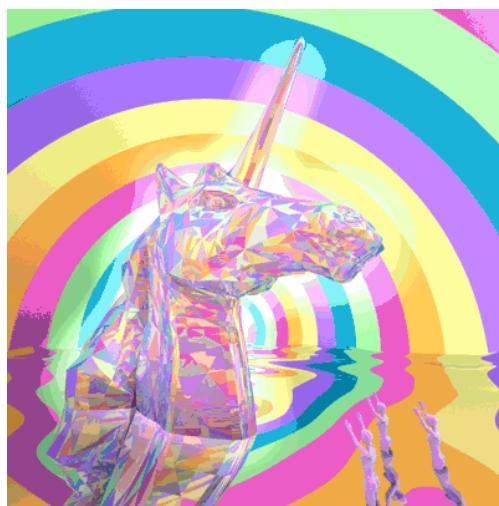
Only the netlify instance was configured securely and included HSTS to prevent MITM SSL stripping attacks.

However, a combination of arp+DNS spoofing allows an attacker to redirect a victim machine (tested on Ubuntu 20.4, Brave, Firefox, Chrome) to an arbitrary site bypassing the HSTS.



## Remote testing

We conducted brief infrastructure tests to understand potential, potentially scalable remote threats by a less privileged attacker. At some point, we discontinued the research due to legal risks. Up until then, no findings were made, i.e., the launchpad's security in this threat model is summarized in the graph below :)



The following table documents testing efforts and summarizes the non-findings:

|                | Manual audit     | hidden site/directory fuzzing | Port probing | http-request smuggling | Server cache poisoning | Metasploit, etc. |
|----------------|------------------|-------------------------------|--------------|------------------------|------------------------|------------------|
| AWS S3         | Not applicable   | No signal                     | No signal    | Limited signal         | discontinued           | untested         |
| mainnet.infura | Not applicable   | No signal                     | untested     | No signal              | untested               | untested         |
| Launchpad      | 1 false positive | No signal                     | No signal    | No signal              | untested               | untested         |

## **cross client issue - libp2p is susceptible to TCP-reset attacks**

### **Location**

<https://github.com/prysmaticlabs/prysm/blob/develop/beacon-chain/p2p/handshake.go#L153>

### **Synopsis**

An attacker can learn the IP addresses and open TCP ports via libp2p. A malicious actor may send large volumes of TCP packets with spoofed source/destination IP and TCP "RST" set to true. Source and destination nodes will disconnect their TCP streams when receiving and processing the spoofed packets **iff** the packet's sequence numbers are in the correct TCP window size.

Therefore, there are three subtypes for this attack to consider;

- a) Blind TCP-reset attacks;

This subtype is essentially  $\geq$  TCP based DOS, where an attacker spams the libp2p network with TCP reset packets with arbitrary sequence numbers.

- b) Educated TCP-reset attacks; an attacker with a data science hat. Interesting entry points for this attack are:

- a newly restarted node (rather low entropy for the attack as low sequence numbers, small amount of connections)
- Established and fully synced nodes with long-lived connections (TCP packets/sequence numbers per connected hour/day/month are observable by connecting to victim nodes through libp2p)

- c) Responsive TCP-reset attacks;

A malicious actor can perform this attack, if it can directly observe the TCP stream, i.e. by an actor with access to LAN, VPS providers, ISPs

# Impact

The naive PoC of a responsive TCP-reset attack had significant impact on all connections of the Prysm beacon node:

The beacon-node will attempt to replace any reset connection with a new one. With a naive, single-threaded implementation of the responsive TCP-reset attack, the connectivity slowly degenerated but was however never brought down to 0. We will further investigate how the victim node's peer score attestation/proposal rates are affected.

Good job though, libp2p discovery!



However, the attack is easy to implement and scale. Furthermore, the attack is inexpensive enough to perform over an extended time on a series of block proposers and attesters. This makes it possible that no new blocks are getting gossipped, which threatens the liveness of the chain and can lead to inactivity penalties.

## Notes & Nitpicks

### Node can force to become Secio “preferred peer”

Location:

<https://github.com/libp2p/specs/blob/master/secio/README.md#determining-roles-and-algorithms>

Secio describes a proof-of-work-like scheme, where two hashes (“oh1 < oh2”) are compared to determine which cipher suites are available for an upcoming handshake. The hash is computed from a concatenation that includes a nonce, which may be arbitrarily chosen by an attacker. Therefore, an attacker can, with high likelihood, become the “preferred peer” in this scheme and force its selection of exchanges, suites, and hashes.

### Prysm - http-web3 provider endpoint accepts http cleartext without warning

That's not a sane default.

### Eth2.0-deposit-CLI insecure default Keystore permissions generated by the

Location:

<https://github.com/ethereum/eth2.0-deposit-cli/>

The default permissions of generated validator keys are 644.

- processes ran by other users on the machine can read the keystore
- processes ran by the same users can read/write the keystore (can be very bad)

Private key permissions should be set to 400, similar to ssh keys.

## Next Steps

### Continuous beacon-fuzzing - Prysm

We conducted some 10<sup>7</sup> brute force TCP fuzzing iterations on prysm (fuzzotron + radamsa). A dedicated beacon fuzzing instance is currently being set up and planned to be available early 2020 for continuous fuzzing of the latest prysm.

### Twist/Invalid Curve attacks on secp256k1

Several parties claimed secp256k1 is not twist secure. Implications on Geth/Libp2p should be understood and if found to be significant, the vulnerabilities should be further investigated.

### Manual auditing - Prysm

For the upcoming two months, a large portion of time is set aside for manual auditing, in particular leveraging @protolamdas attacknets and rumor toolstack. Additionally, noise/secio protocol downgrades seem to be interesting leads

### Social Engineering, Spear phishing - Ethereum.org

In particular, spear phishing is still one of the most prevalent threats to most organizations. Therefore, we recommend to conduct a targeted, Gsuite based (2fa-) credential harvesting spear-phishing campaign to collect data on how effective such an attack would be and further harden contributors' senses for this particular threat.